

Reference

This is a DataCamp course

Preparing data

Reading DataFrames from multiple files

```
In [1]: import pandas as pd

bronze = pd.read_csv('Bronze.csv')
silver = pd.read_csv('Silver.csv')
gold = pd.read_csv('Gold.csv')
print(gold.head())
```

	NOC	Country	Total
0	USA	United States	2088.0
1	URS	Soviet Union	838.0
2	GBR	United Kingdom	498.0
3	FRA	France	378.0
4	GER	Germany	407.0

Combining DataFrames from multiple data files

```
In [ ]: import pandas as pd
medals = gold.copy()
new_labels = ['NOC', 'Country', 'Gold']
medals.columns = new_labels
medals['Silver'] = silver['Total']
medals['Bronze'] = bronze['Total']
print(medals.head())
```

Sorting DataFrame with the Index & columns

The following is similar to 'order by' in SQL

```
In [2]: import pandas as pd

weatherDic = { 'Max TemperatureF':[68, 60, 68, 84, 88], 'Month':['Jan', 'Feb', 'Mar', 'Apr', 'May']}
weather1 = pd.DataFrame(weatherDic)
weather1 = weather1.set_index('Month')
weather2 = weather1.sort_index()
weather3 = weather1.sort_index(ascending=False)
weather4 = weather1.sort_values('Max TemperatureF')
```

Reindexing DataFrame from a list

Be familiar with `.set_index()` and `.reindex()`

```

In [3]: import pandas as pd

year = ['Jan',
        'Feb',
        'Mar',
        'Apr',
        'May',
        'Jun',
        'Jul',
        'Aug',
        'Sep',
        'Oct',
        'Nov',
        'Dec']

weatherDic = { 'Mean TemperatureF':[61.956044, 32.133333, 68.934783, 43.434783], 'Month':['Apr', 'Jan', 'Jul', 'Oct']
weather1 = pd.DataFrame(weatherDic)
weather1 = weather1.set_index('Month')

weather2 = weather1.reindex(year)

print(weather2)

weather3 = weather1.reindex(year).ffill()

print(weather3)

```

Month	Mean TemperatureF
Jan	32.133333
Feb	NaN
Mar	NaN
Apr	61.956044
May	NaN
Jun	NaN
Jul	68.934783
Aug	NaN
Sep	NaN
Oct	43.434783
Nov	NaN
Dec	NaN

Mean TemperatureF	
Month	
Jan	32.133333
Feb	32.133333
Mar	32.133333
Apr	61.956044
May	61.956044
Jun	61.956044
Jul	68.934783
Aug	68.934783
Sep	68.934783
Oct	43.434783
Nov	43.434783
Dec	43.434783

Reindexing using another DataFrame Index

Note the multi-Index

```
In [4]: import pandas as pd

names_1981 = pd.read_csv("names1981.csv")
names_1881 = pd.read_csv("names1881.csv")

column_names = ['name', 'gender', 'count']
names_1981.columns = column_names
names_1881.columns = column_names

names_1981 = names_1981.set_index(['name', 'gender'])
names_1881 = names_1881.set_index(['name', 'gender'])
print(names_1881.head())
print(names_1981.head())

print(names_1881.shape)
print(names_1981.shape)

common_names = names_1981.reindex(names_1881.index)
print(common_names.shape)
common_names = common_names.dropna()
print(common_names.shape)
```

		count
name	gender	
Anna	F	2698
Emma	F	2034
Elizabeth	F	1852
Margaret	F	1658
Minnie	F	1653

		count
name	gender	
Jessica	F	42519
Amanda	F	34370
Sarah	F	28162
Melissa	F	28003
Amy	F	20337
(1934,	1)	
(19454,	1)	
(1934,	1)	
(1586,	1)	

Adding unaligned DataFrames

If you were to add the following two DataFrames by executing the command `total = january + february`, how many rows would the resulting DataFrame have?

january

Units

Company

Acme Corporation 19

Hooli 17

Initech 20

Mediacore 10

Streeplex 13

february

Units

Company

Acme Corporation 15

Hooli 3

Mediacore 13

Vandelay Inc 25

Answer:

january and february both consist of the sales of the Companies Acme Corporation, Hooli, and Mediacore. january has the additional two companies Initech and Streeplex, while february has the additional company Vandelay Inc. Together, they consist of the sales of 6 unique companies, and so total would have 6. **So this is like the full-join in SQL to some extent.**

Broadcasting in arithmetic formulas

```
In [39]: import pandas as pd
weather = pd.read_csv('pittsburgh2013.csv', index_col = 'Date')
temps_f = weather[['Min TemperatureF', 'Mean TemperatureF', 'Max TemperatureF']]
temps_c = (temps_f - 32) * 5/9
temps_c.columns = temps_c.columns.str.replace('F', 'C')
# Be aware of the .replace().

print(temps_c.head())
```

	Min TemperatureC	Mean TemperatureC	Max TemperatureC
Date			
2013-1-1	-6.111111	-2.222222	0.000000
2013-1-2	-8.333333	-6.111111	-3.888889
2013-1-3	-8.888889	-4.444444	0.000000
2013-1-4	-2.777778	-2.222222	-1.111111
2013-1-5	-3.888889	-1.111111	1.111111

Computing percentage growth of GDP

Compute the percentage growth of the resampled DataFrame yearly with `.pct_change() * 100`, which is defined on either Pandas Series or DataFrame.

```
In [7]: import pandas as pd
gdp = pd.read_csv('gdp_usa.csv', index_col='DATE', parse_dates=True)
post2008 = gdp['2008:']
yearly = post2008.resample('A').last() # The original is quartly data and thus resample by year.
yearly['growth'] = yearly.pct_change() * 100
print(yearly)
```

	VALUE	growth
DATE		
2008-12-31	14549.9	NaN
2009-12-31	14566.5	0.114090
2010-12-31	15230.2	4.556345
2011-12-31	15785.3	3.644732
2012-12-31	16297.3	3.243524
2013-12-31	16999.9	4.311144
2014-12-31	17692.2	4.072377
2015-12-31	18222.8	2.999062
2016-12-31	18436.5	1.172707

Converting currency of stocks

```
In [15]: import pandas as pd
sp500 = pd.read_csv('sp500.csv', index_col='Date', parse_dates=True)
exchange = pd.read_csv('exchange.csv', index_col='Date', parse_dates=True)
dollars = sp500[['Open', 'Close']]
pounds = dollars.multiply(exchange['GBP/USD'], axis=0)
print(pounds.head())
```

	Open	Close
Date		
2015-01-02	1340.364425	1339.908750
2015-01-05	1348.616555	1326.389506
2015-01-06	1332.515980	1319.639876
2015-01-07	1330.562125	1344.063112
2015-01-08	1343.268811	1364.126161

Concatenating data

Perform database-style operations to combine DataFrames: appending and concatenating DataFrames.

Appending pandas Series

```
In [17]: import pandas as pd
jan = pd.read_csv('sales-jan-2015.csv', index_col='Date', parse_dates=True)
feb = pd.read_csv('sales-feb-2015.csv', index_col='Date', parse_dates=True)
mar = pd.read_csv('sales-mar-2015.csv', index_col='Date', parse_dates=True)

jan_units = jan['Units']
feb_units = feb['Units']
mar_units = mar['Units']

quarter1 = jan_units.append(feb_units).append(mar_units)
print(jan_units.shape)
print(feb_units.shape)
print(mar_units.shape)
print(quarter1.shape)

print(quarter1.loc['jan 27, 2015':'feb 2, 2015'])
print(quarter1.loc['feb 26, 2015':'mar 7, 2015'])
print(quarter1.sum())
```

```
(20,)
(20,)
(20,)
(60,)
Date
2015-01-27 07:11:55    18
2015-02-02 08:33:01     3
2015-02-02 20:54:49     9
Name: Units, dtype: int64
Date
2015-02-26 08:57:45     4
2015-02-26 08:58:51     1
2015-03-06 10:11:45    17
2015-03-06 02:03:56    17
Name: Units, dtype: int64
642
```

****The above appending DataFrame or Series is like union/union all of SQL set clause, which operate on rows?**

Concatenating pandas Series along row axis

Having learned how to append Series, now learn how to achieve the same result by concatenating Series instead. Then **what is the difference between `pd.concat()` and pandas' `.append()` method**. One way to think of the difference is that `.append()` is a specific case of a concatenation, while `pd.concat()` gives you more flexibility, as you'll see in later exercises.

```
In [47]: units = []

for month in [jan, feb, mar]:
    units.append(month['Units'])

quarter1 = pd.concat(units, axis='rows')

print(quarter1.loc['jan 27, 2015':'feb 2, 2015'])
print(quarter1.loc['feb 26, 2015':'mar 7, 2015'])
```

```
Date
2015-01-27 07:11:55    18
2015-02-02 08:33:01     3
2015-02-02 20:54:49     9
Name: Units, dtype: int64
Date
2015-02-26 08:57:45     4
2015-02-26 08:58:51     1
2015-03-06 10:11:45    17
2015-03-06 02:03:56    17
Name: Units, dtype: int64
```

Appending DataFrames with ignore_index

DataFrames `names_1981` and `names_1881` are loaded without specifying an Index column (so the default Indexes for both are **RangeIndexes**).

Use the DataFrame `.append()` method to make a DataFrame `combined_names`. To distinguish rows from the original two DataFrames, you'll add a 'year' column to each with the year (1881 or 1981 in this case). In addition, Specify `ignore_index=True` so that the index values are not used along the concatenation axis. The resulting axis will instead be labeled 0, 1, ..., n-1, which is useful if you are concatenating objects where the concatenation axis does not have meaningful indexing information.


```
In [64]: import pandas as pd
names_1981 = pd.read_csv("names1981.csv")
names_1881 = pd.read_csv("names1881.csv")

columnsList = ['name', 'gender', 'count']
names_1981.columns = columnsList
names_1881.columns = columnsList

# Add 'year' column to names_1881 & names_1981
names_1881['year'] = 1881
names_1981['year'] = 1981

# Append names_1981 after names_1881 with ignore_index=True: combined_names
combined_names = names_1881.append(names_1981, ignore_index=True)
#This will use the default RangeIndex
print(combined_names.index)
# This will give a different index. May be compile the original together?
combined_names1 = names_1881.append(names_1981)
print(combined_names1.index)

# Print shapes of names_1981, names_1881, and combined_names
print(names_1981.shape)
print(names_1881.shape)
print(combined_names.shape)

# Print all rows that contain the name 'Morgan'
print(combined_names.loc[combined_names['name']=='Morgan'])
```

```
RangeIndex(start=0, stop=21388, step=1)
Int64Index([ 0, 1, 2, 3, 4, 5, 6, 7, 8,
            9,
            ...,
            19444, 19445, 19446, 19447, 19448, 19449, 19450, 19451, 19452,
            19453],
            dtype='int64', length=21388)
(19454, 4)
(1934, 4)
(21388, 4)
   name gender  count  year
1282  Morgan    M     23  1881
2094  Morgan    F    1769  1981
14388  Morgan    M     766  1981
```

Concatenating pandas DataFrames along column axis

The function `pd.concat()` can concatenate DataFrames horizontally as well as vertically (vertical is the default). To make the DataFrames stack horizontally, you have to specify the keyword argument `axis=1` or `axis='columns'`. **Here we know that unlike set clauses in SQL, `pd.concat()` can combine DataFrames both horizontally and vertically.**

```
In [73]: import pandas as pd
weather_max = pd.DataFrame({'Month': ['Jan', 'Apr', 'Jul', 'Oct'], 'Max TemperatureF': [68, 89, 91, 84]})
weather_max = weather_max.set_index('Month')
print(weather_max)
print('-----')
weather_mean = pd.DataFrame({'Month': ['Apr', 'Aug', 'Dec', 'Feb', 'Jan', 'Jul', 'Jun', 'Mar', 'May', 'Nov', 'Oct', 'Sep'],
                             'Mean TemperatureF': [53.100000, 70.000000, 34.935484, 28.714286, 32.354839, 72.870968, 70.133333,
                                                     35.000000, 62.612903, 39.800000, 55.451613, 63.766667]})
weather_mean = weather_mean.set_index('Month')
print(weather_mean)
print('-----')

weather = pd.concat([weather_max, weather_mean], axis=1)

print(weather)
```

```

      Max TemperatureF
Month
Jan                68
Apr                89
Jul                91
Oct                84
-----
```

```

      Mean TemperatureF
Month
Apr          53.100000
Aug          70.000000
Dec          34.935484
Feb          28.714286
Jan          32.354839
Jul          72.870968
Jun          70.133333
Mar          35.000000
May          62.612903
Nov          39.800000
Oct          55.451613
Sep          63.766667
-----
```

```

      Max TemperatureF  Mean TemperatureF
Apr                89.0          53.100000
Aug                NaN          70.000000
Dec                NaN          34.935484
```

Feb	NaN	28.714286
Jan	68.0	32.354839
Jul	91.0	72.870968
Jun	NaN	70.133333
Mar	NaN	35.000000
May	NaN	62.612903
Nov	NaN	39.800000
Oct	84.0	55.451613
Sep	NaN	63.766667

Reading multiple files to build a DataFrame

```
In [2]: import pandas as pd
medal_types = ['bronze', 'silver', 'gold']
medals = []
for medal in medal_types:
    file_name = "%s_top5.csv" % medal
    columns = ['Country', medal]
    medal_df = pd.read_csv(file_name, header=0, index_col='Country', names=columns)
    medals.append(medal_df)

medals = pd.concat(medals, axis='columns')

print(medals)
```

	bronze	silver	gold
France	475.0	461.0	NaN
Germany	454.0	NaN	407.0
Italy	NaN	394.0	460.0
Soviet Union	584.0	627.0	838.0
United Kingdom	505.0	591.0	498.0
United States	1052.0	1195.0	2088.0

Concatenating vertically to get MultiIndexed rows

When stacking a sequence of DataFrames vertically, it is sometimes desirable to construct a MultiIndex to indicate the DataFrame from which each row originated.

```

In [6]: medal_types = ['bronze', 'silver', 'gold']
        medals = []

        for medal in medal_types:

            file_name = "%s_top5.csv" % medal

            # Read file_name into a DataFrame: medal_df
            medal_df = pd.read_csv(file_name, index_col='Country')
            #print(medal_df.index)
            # Append medal_df to medals
            medals.append(medal_df)

        print(medals[0])
        print('-----')
        print(medals[1])
        print('-----')
        print(medals[2])
        print('-----')

        # Concatenate medals: medals

        # medals = pd.concat(medals) #Uncomment the following three and comment the later, and see what happens if no key
        # # Print medals
        # print(medals)
        # print('-----')

        medals = pd.concat(medals, keys=['bronze', 'silver', 'gold'])
        print(medals)
        print('-----')
        print(medals.index)

```

	Total
Country	
United States	1052.0
Soviet Union	584.0
United Kingdom	505.0
France	475.0
Germany	454.0

	Total
Country	
United States	1195.0
Soviet Union	627.0
United Kingdom	591.0
France	461.0
Italy	394.0

	Total
Country	
United States	2088.0
Soviet Union	838.0
United Kingdom	498.0
Italy	460.0
Germany	407.0

	Country	Total
bronze	United States	1052.0
	Soviet Union	584.0
	United Kingdom	505.0
	France	475.0
	Germany	454.0
silver	United States	1195.0
	Soviet Union	627.0
	United Kingdom	591.0
	France	461.0
	Italy	394.0
gold	United States	2088.0
	Soviet Union	838.0
	United Kingdom	498.0
	Italy	460.0
	Germany	407.0

```
MultiIndex(levels=[['bronze', 'silver', 'gold'], ['France', 'Germany', 'Italy', 'Soviet Union', 'United Kingdom', 'United States']],
            labels=[[0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 2, 2, 2, 2, 2], [5, 3, 4, 0, 1, 5, 3, 4, 0, 2, 5, 3, 4, 2, 1]],
            names=[None, 'Country'])
```

Total

Country

United States 1052.0 Soviet Union 584.0 United Kingdom 505.0 France 475.0 Germany 454.0 United States 1195.0 Soviet Union 627.0
United Kingdom 591.0 France 461.0 Italy 394.0 United States 2088.0 Soviet Union 838.0 United Kingdom 498.0 Italy 460.0 Germany
407.0

Slicing MultiIndexed DataFrames

This exercise picks up where the last ended (again using The Guardian's Olympic medal dataset).

You are provided with the MultiIndexed DataFrame as produced at the end of the preceding exercise. Your task is to sort the DataFrame and to use the `pd.IndexSlice` to extract specific slices. Check out this exercise from [Manipulating DataFrames with pandas](#) to refresh your memory on how to deal with MultiIndexed DataFrames.

`pandas` has been imported for you as `pd` and the DataFrame `medals` is already in your namespace.

INSTRUCTIONS

Create a new DataFrame `medals_sorted` with the entries of `medals` sorted. Use `.sort_index(level=0)` to ensure the Index is sorted suitably. Print the number of bronze medals won by Germany and all of the silver medal data. This has been done for you. Create an alias for `pd.IndexSlice` called `idx`. **A slicer `pd.IndexSlice` is required when slicing on the inner level of a MultiIndex.** Slice all the data on medals won by the United Kingdom. To do this, use the `.loc[]` accessor with `idx[:, 'United Kingdom']`, `..`

```

In [101]: print(medals)
print('-----')
# Sort the entries of medals
medals_sorted = medals.sort_index(level=0) #print(medals.index) can show what are level=0 index.

print(medals_sorted)
print('-----')

# Print the number of Bronze medals won by Germany
print(medals_sorted.loc[('bronze', 'Germany')]) #Note how to locate multi-index with loc.
print('-----')

# Print data about silver medals
print(medals_sorted.loc['silver'])
print('-----')

# Create alias for pd.IndexSlice: idx
idx = pd.IndexSlice

# Print all the data on medals won by the United Kingdom
print(medals_sorted.loc[idx[:, 'United Kingdom'], :])
print('-----')

```

	Country	Total
bronze	United States	1052.0
	Soviet Union	584.0
	United Kingdom	505.0
	France	475.0
	Germany	454.0
silver	United States	1195.0
	Soviet Union	627.0
	United Kingdom	591.0
	France	461.0
	Italy	394.0
gold	United States	2088.0
	Soviet Union	838.0
	United Kingdom	498.0
	Italy	460.0
	Germany	407.0

	Country	Total
bronze	France	475.0
	Germany	454.0
	Soviet Union	584.0
	United Kingdom	505.0
	United States	1052.0
gold	Germany	407.0
	Italy	460.0
	Soviet Union	838.0
	United Kingdom	498.0
	United States	2088.0
silver	France	461.0
	Italy	394.0
	Soviet Union	627.0
	United Kingdom	591.0
	United States	1195.0

Total	454.0
-------	-------

Name: (bronze, Germany), dtype: float64

	Country	Total
	France	461.0
	Italy	394.0
	Soviet Union	627.0
	United Kingdom	591.0
	United States	1195.0

	Country	Total
bronze	United Kingdom	505.0
gold	United Kingdom	498.0
silver	United Kingdom	591.0

Concatenating horizontally to get MultiIndexed columns

It is also possible to construct a DataFrame with hierarchically indexed columns. For this exercise, you'll start with pandas imported and a list of three DataFrames called `dataframes`. All three DataFrames contain 'Company', 'Product', and 'Units' columns with a 'Date' column as the index pertaining to sales transactions during the month of February, 2015. The first DataFrame describes Hardware transactions, the second describes Software transactions, and the third, Service transactions.

Your task is to concatenate the DataFrames horizontally and to create a MultiIndex on the columns. From there, you can summarize the resulting DataFrame and slice some information from it.

INSTRUCTIONS

Construct a new DataFrame february with MultiIndexed columns by concatenating the list dataframes. Use axis=1 to stack the DataFrames horizontally and the keyword argument keys=['Hardware', 'Software', 'Service'] to construct a hierarchical Index from each DataFrame. Print summary information from the new DataFrame february using the .info() method. This has been done for you. Create an alias called idx for pd.IndexSlice. Extract a slice called slice_2_8 from february (using .loc[] & idx) that comprises rows between Feb. 2, 2015 to Feb. 8, 2015 from columns under 'Company'. Print the slice_2_8. This has been done for you, so hit 'Submit Answer' to see the sliced data!

```

In [110]: import pandas as pd

hardware = pd.read_csv('feb-sales-Hardware.csv', index_col='Date', parse_dates=True)
software = pd.read_csv('feb-sales-Software.csv', index_col='Date', parse_dates=True)
service = pd.read_csv('feb-sales-Service.csv', index_col='Date', parse_dates=True)

dataframes = [hardware, software, service]
print(dataframes[0])
print('-----')
print(dataframes[1])
print('-----')

print(dataframes[2])
print('-----')

# my code above

# # Concatenate dataframes: february
# february = pd.concat(dataframes, axis=1) #Uncomment the following three and comment the three later and compare
# print(february)
# print('-----')

february = pd.concat(dataframes, keys=['Hardware', 'Software', 'Service'], axis=1)
print(february)
print('-----')

# Print february.info()
print(february.info())
print('-----')

# Assign pd.IndexSlice: idx
idx = pd.IndexSlice

# Create the slice: slice_2_8
slice_2_8 = february.loc['2015-2-2':'2015-2-8', idx[:, 'Company']]

# Print slice_2_8
print(slice_2_8)

```

Date	Company	Product	Units
------	---------	---------	-------

2015-02-04	21:52:45	Acme Coporation	Hardware	14
2015-02-07	22:58:10	Acme Coporation	Hardware	1
2015-02-19	10:59:33	Mediacore	Hardware	16
2015-02-02	20:54:49	Mediacore	Hardware	9
2015-02-21	20:41:47	Hooli	Hardware	3

Date		Company	Product	Units
2015-02-16	12:09:19	Hooli	Software	10
2015-02-03	14:14:18	Initech	Software	13
2015-02-02	08:33:01	Hooli	Software	3
2015-02-05	01:53:06	Acme Coporation	Software	19
2015-02-11	20:03:08	Initech	Software	7
2015-02-09	13:09:55	Mediacore	Software	7
2015-02-11	22:50:44	Hooli	Software	4
2015-02-04	15:36:29	Streeplex	Software	13
2015-02-21	05:01:26	Mediacore	Software	3

Date		Company	Product	Units
2015-02-26	08:57:45	Streeplex	Service	4
2015-02-25	00:29:00	Initech	Service	10
2015-02-09	08:57:30	Streeplex	Service	19
2015-02-26	08:58:51	Streeplex	Service	1
2015-02-05	22:05:03	Hooli	Service	10
2015-02-19	16:02:58	Mediacore	Service	10

Date		Hardware Company	Product	Units	Software Company \
2015-02-02	08:33:01	NaN	NaN	NaN	Hooli
2015-02-02	20:54:49	Mediacore	Hardware	9.0	NaN
2015-02-03	14:14:18	NaN	NaN	NaN	Initech
2015-02-04	15:36:29	NaN	NaN	NaN	Streeplex
2015-02-04	21:52:45	Acme Coporation	Hardware	14.0	NaN
2015-02-05	01:53:06	NaN	NaN	NaN	Acme Coporation
2015-02-05	22:05:03	NaN	NaN	NaN	NaN
2015-02-07	22:58:10	Acme Coporation	Hardware	1.0	NaN
2015-02-09	08:57:30	NaN	NaN	NaN	NaN
2015-02-09	13:09:55	NaN	NaN	NaN	Mediacore
2015-02-11	20:03:08	NaN	NaN	NaN	Initech
2015-02-11	22:50:44	NaN	NaN	NaN	Hooli
2015-02-16	12:09:19	NaN	NaN	NaN	Hooli

2015-02-19 10:59:33	Mediacore	Hardware	16.0	NaN
2015-02-19 16:02:58	NaN	NaN	NaN	NaN
2015-02-21 05:01:26	NaN	NaN	NaN	Mediacore
2015-02-21 20:41:47	Hooli	Hardware	3.0	NaN
2015-02-25 00:29:00	NaN	NaN	NaN	NaN
2015-02-26 08:57:45	NaN	NaN	NaN	NaN
2015-02-26 08:58:51	NaN	NaN	NaN	NaN

Date	Product	Units	Service Company	Product	Units
2015-02-02 08:33:01	Software	3.0	NaN	NaN	NaN
2015-02-02 20:54:49	NaN	NaN	NaN	NaN	NaN
2015-02-03 14:14:18	Software	13.0	NaN	NaN	NaN
2015-02-04 15:36:29	Software	13.0	NaN	NaN	NaN
2015-02-04 21:52:45	NaN	NaN	NaN	NaN	NaN
2015-02-05 01:53:06	Software	19.0	NaN	NaN	NaN
2015-02-05 22:05:03	NaN	NaN	Hooli	Service	10.0
2015-02-07 22:58:10	NaN	NaN	NaN	NaN	NaN
2015-02-09 08:57:30	NaN	NaN	Streeplex	Service	19.0
2015-02-09 13:09:55	Software	7.0	NaN	NaN	NaN
2015-02-11 20:03:08	Software	7.0	NaN	NaN	NaN
2015-02-11 22:50:44	Software	4.0	NaN	NaN	NaN
2015-02-16 12:09:19	Software	10.0	NaN	NaN	NaN
2015-02-19 10:59:33	NaN	NaN	NaN	NaN	NaN
2015-02-19 16:02:58	NaN	NaN	Mediacore	Service	10.0
2015-02-21 05:01:26	Software	3.0	NaN	NaN	NaN
2015-02-21 20:41:47	NaN	NaN	NaN	NaN	NaN
2015-02-25 00:29:00	NaN	NaN	Initech	Service	10.0
2015-02-26 08:57:45	NaN	NaN	Streeplex	Service	4.0
2015-02-26 08:58:51	NaN	NaN	Streeplex	Service	1.0

```

-----
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 20 entries, 2015-02-02 08:33:01 to 2015-02-26 08:58:51
Data columns (total 9 columns):
(Hardware, Company)    5 non-null object
(Hardware, Product)    5 non-null object
(Hardware, Units)      5 non-null float64
(Software, Company)    9 non-null object
(Software, Product)    9 non-null object
(Software, Units)      9 non-null float64
(Service, Company)     6 non-null object
(Service, Product)     6 non-null object

```



```
(Service, Units)      6 non-null float64
dtypes: float64(3), object(6)
memory usage: 1.6+ KB
None
```

```
-----
                                Hardware      Software Service
                                Company      Company Company
Date
2015-02-02 08:33:01              NaN          Hooli      NaN
2015-02-02 20:54:49          Mediacore          NaN      NaN
2015-02-03 14:14:18              NaN          Initech      NaN
2015-02-04 15:36:29              NaN        Streeplex      NaN
2015-02-04 21:52:45  Acme Coporation          NaN      NaN
2015-02-05 01:53:06              NaN  Acme Coporation      NaN
2015-02-05 22:05:03              NaN          NaN      Hooli
2015-02-07 22:58:10  Acme Coporation          NaN      NaN
```

Concatenating DataFrames from a dict

You're now going to revisit the sales data you worked with earlier in the chapter. Three DataFrames `jan`, `feb`, and `mar` have been pre-loaded for you. Your task is to aggregate the sum of all sales over the 'Company' column into a single DataFrame. You'll do this by constructing a dictionary of these DataFrames and then concatenating them.

INSTRUCTIONS

Create a list called `month_list` consisting of the tuples `('january', jan)`, `('february', feb)`, and `('march', mar)`. Create an empty dictionary called `month_dict`. Inside the for loop: Group `month_data` by 'Company' and use `.sum()` to aggregate. Construct a new DataFrame called `sales` by concatenating the DataFrames stored in `month_dict`. Create an alias for `pd.IndexSlice` and print all sales by 'Mediacore'. This has been done for you, so hit 'Submit Answer' to see the result!

```
In [10]: # Import pandas
import pandas as pd

# Load 'sales-jan-2015.csv' into a DataFrame: jan
jan = pd.read_csv('sales-jan-2015.csv', index_col='Date', parse_dates=True)

# Load 'sales-feb-2015.csv' into a DataFrame: feb
feb = pd.read_csv('sales-feb-2015.csv', index_col='Date', parse_dates=True)

# Load 'sales-mar-2015.csv' into a DataFrame: mar
mar = pd.read_csv('sales-mar-2015.csv', index_col='Date', parse_dates=True)

#my code above

# Make the list of tuples: month_list
month_list = [('january', jan), ('february', feb), ('march', mar)]

# Create an empty dictionary: month_dict
month_dict = {}

for month_name, month_data in month_list:

    # Group month_data: month_dict[month_name]
    month_dict[month_name] = month_data.groupby('Company').sum()
    print(month_name)
    print('-----')
    print(month_data)
    print('-----')
    print(month_dict[month_name]) ##Check why there are fewer columns after group by. Seems not as before.
    print('-----')

# Concatenate data in month_dict: sales
sales = pd.concat(month_dict)

# Print sales
print(sales)

# Print all sales by Mediacore
idx = pd.IndexSlice
print(sales.loc[idx[:, 'Mediacore'], :])
```

january

Date	Company	Product	Units
2015-01-21 19:13:21	Streeplex	Hardware	11
2015-01-09 05:23:51	Streeplex	Service	8
2015-01-06 17:19:34	Initech	Hardware	17
2015-01-02 09:51:06	Hooli	Hardware	16
2015-01-11 14:51:02	Hooli	Hardware	11
2015-01-01 07:31:20	Acme Coporation	Software	18
2015-01-24 08:01:16	Initech	Software	1
2015-01-25 15:40:07	Initech	Service	6
2015-01-13 05:36:12	Hooli	Service	7
2015-01-03 18:00:19	Hooli	Service	19
2015-01-16 00:33:47	Hooli	Hardware	17
2015-01-16 07:21:12	Initech	Service	13
2015-01-20 19:49:24	Acme Coporation	Hardware	12
2015-01-26 01:50:25	Acme Coporation	Software	14
2015-01-15 02:38:25	Acme Coporation	Service	16
2015-01-06 12:47:37	Acme Coporation	Service	16

There is a question about group by above

Concatenating DataFrames with inner join

Here, you'll continue working with DataFrames compiled from The Guardian's Olympic medal dataset.

The DataFrames bronze, silver, and gold have been pre-loaded for you.

Your task is to compute an inner join.

INSTRUCTIONS

Construct a list of DataFrames called medal_list with entries bronze, silver, and gold. Concatenate medal_list horizontally with an inner join to create medals. Use the keyword argument keys=['bronze', 'silver', 'gold'] to yield suitable hierarchical indexing. Use axis=1 to get horizontal concatenation. Use join='inner' to keep only rows that share common index labels. Print the new DataFrame medals.

```
In [20]: import pandas as pd

bronze = pd.DataFrame({'Country': ['United States', 'Soviet Union', 'United Kingdom', 'France', 'Germany'],
                        'Total': [1052.0, 584.0, 505.0, 475.0, 454.0]})
bronze = bronze.set_index('Country')

silver = pd.DataFrame({'Country': ['United States', 'Soviet Union', 'United Kingdom', 'France', 'Italy'],
                        'Total': [1195.0, 627.0, 591.0, 461.0, 394.0]})
silver = silver.set_index('Country')

gold = pd.DataFrame({'Country': ['United States', 'Soviet Union', 'United Kingdom', 'Italy', 'Germany'],
                      'Total': [2088.0, 838.0, 498.0, 460.0, 407.0]})
gold = gold.set_index('Country')

#my code above

# Create the list of DataFrames: medal_list
medal_list = [bronze, silver, gold]

# Concatenate medal_list horizontally using an inner join: medals
medals = pd.concat(medal_list, keys=['bronze', 'silver', 'gold'], axis=1, join='inner')

# Print medals
print(medals)
```

	bronze	silver	gold
Country	Total	Total	Total
United States	1052.0	1195.0	2088.0
Soviet Union	584.0	627.0	838.0
United Kingdom	505.0	591.0	498.0

Resampling & concatenating DataFrames with inner join

In this exercise, you'll compare the historical 10-year GDP (Gross Domestic Product) growth in the US and in China. The data for the US starts in 1947 and is recorded quarterly; by contrast, the data for China starts in 1961 and is recorded annually.

You'll need to use a combination of resampling and an inner join to align the index labels. You'll need an appropriate offset alias for resampling, and the method `.resample()` must be chained with some kind of aggregation method (`.pct_change()` and `.last()` in this case).

pandas has been imported as `pd`, and the DataFrames `china` and `us` have been pre-loaded, with the output of `china.head()` and `us.head()` printed in the IPython Shell.

INSTRUCTIONS

Make a new DataFrame `china_annual` by resampling the DataFrame `china` with `.resample('A')` (i.e., with annual frequency) and chaining two method calls: Chain `.pct_change(10)` as an aggregation method to compute the percentage change with an offset of ten years. Chain `.dropna()` to eliminate rows containing null values. Make a new DataFrame `us_annual` by resampling the DataFrame `us` exactly as you resampled `china`. Concatenate `china_annual` and `us_annual` to construct a DataFrame called `gdp`. Use `join='inner'` to perform an inner join and use `axis=1` to concatenate horizontally. Print the result of resampling `gdp` every decade (i.e., using `.resample('10A')`) and aggregating with the method `.last()`. This has been done for you, so hit 'Submit Answer' to see the result!

I need make it clear about the usage of 'A' and '10A' etc.

```
In [39]: import pandas as pd
china = pd.read_csv("gdp_china.csv", index_col = "Year", parse_dates = True )
china.columns = ['China'] #### This needs []
china.index.name = 'Year' #### This does not need []
print(china.head())
us = pd.read_csv("gdp_usa.csv", index_col = "DATE", parse_dates = True)
us.columns = ['US']
us.index.name = 'Year'
print(us.head())

#my code above

# Resample and tidy china: china_annual
china_annual = china.resample('A').mean().pct_change(10).dropna()

# Resample and tidy us: us_annual
us_annual = us.resample('A').mean().pct_change(10).dropna()

# Concatenate china_annual and us_annual: gdp
gdp = pd.concat([china_annual, us_annual], axis=1, join='inner')

print(gdp.head())

# Resample gdp and print
print(gdp.resample('10A').last())
```

China	
Year	
1960-01-01	59.184116
1961-01-01	49.557050
1962-01-01	46.685179
1963-01-01	50.097303
1964-01-01	59.062255

US	
Year	
1947-01-01	243.1
1947-04-01	246.3
1947-07-01	250.1
1947-10-01	260.3
1948-01-01	266.2

	China	US
Year		

1970-12-31	0.546128	0.980397
1971-12-31	0.988860	1.073188
1972-12-31	1.402472	1.119273
1973-12-31	1.730085	1.237090
1974-12-31	1.408556	1.258503
	China	US

Year

1970-12-31	0.546128	0.980397
1980-12-31	1.072537	1.660540
1990-12-31	0.892820	1.088953
2000-12-31	2.357522	0.719980
2010-12-31	4.011081	0.455009
2020-12-31	3.789936	0.377506

In the above example, if `mean()` is deleted, it still will call `mean()` by default and will give warnings.

Merging data

Here, you'll learn all about merging pandas DataFrames. You'll explore different techniques for merging, and learn about left joins, right joins, inner joins, and outer joins, as well as when to use which. You'll also learn about ordered merging, which is useful when you want to merge DataFrames whose columns have natural orderings, like date-time columns.

First clarify the difference between last and this chapters: Concatenating and merging.

Merging company DataFrames

Suppose your company has operations in several different cities under several different managers. The DataFrames `revenue` and `managers` contain partial information related to the company. That is, the rows of the city columns don't quite match in `revenue` and `managers` (the Mendocino branch has no revenue yet since it just opened and the manager of Springfield branch recently left the company).

```
city    revenue
```

```
0 Austin 100 1 Denver 83 2 Springfield 4
```

```
city    manager
```

```
0 Austin Charlers 1 Denver Joel 2 Mendocino Brett
```

The DataFrames have been printed in the IPython Shell. If you were to run the command `combined = pd.merge(revenue, managers, on='city')`, how many rows would combined have?

INSTRUCTIONS

Possible Answers 0 rows. press 1 2 rows. Answer press 2 3 rows. press 3 4 rows. press 4

Remember, the default strategy for `pd.merge()` is an inner join. However, contrast here and other places where it gives full-outer-join like results although it is also called inner join. I must summarize this to distinguish.

Merging on a specific column

This exercise follows on the last one with the DataFrames `revenue` and `managers` for your company. You expect your company to grow and, eventually, to operate in cities with the same name on different states. As such, you decide that every branch should have a numerical branch identifier. Thus, you add a `branch_id` column to both DataFrames. Moreover, new cities have been added to both the `revenue` and `managers` DataFrames as well. `pandas` has been imported as `pd` and both DataFrames are available in your namespace.

At present, there should be a 1-to-1 relationship between the `city` and `branch_id` fields. In that case, the result of a merge on the `city` columns ought to give you the same output as a merge on the `branch_id` columns. Do they? Can you spot an ambiguity in one of the DataFrames?

INSTRUCTIONS

Using `pd.merge()`, merge the DataFrames `revenue` and `managers` on the `'city'` column of each. Store the result as `merge_by_city`. Print the DataFrame `merge_by_city`. This has been done for you. Merge the DataFrames `revenue` and `managers` on the `'branch_id'` column of each. Store the result as `merge_by_id`. Print the DataFrame `merge_by_id`. This has been done for you, so hit 'Submit Answer' to see the result!


```
In [4]: import pandas as pd

revenue = pd.DataFrame({'branch_id':[10,20,30,47], 'city':['Austin', 'Denver', 'Springfield', 'Mendocino'],
                        'revenue':[100,83,4,200]})
print('revenue')
print(revenue)

managers = pd.DataFrame({'branch_id':[10,20,47,31], 'city':['Austin', 'Denver', 'Mendocino', 'Springfield'],
                        'manager':['Charles', 'Joel', 'Brett', 'Sally']})
print('managers')
print(managers)

# Merge revenue with managers on 'city': merge_by_city
merge_by_city = pd.merge(revenue, managers, on='city')

# Print merge_by_city
print(merge_by_city)

# Merge revenue with managers on 'branch_id': merge_by_id
merge_by_id = pd.merge(revenue, managers, on='branch_id')

# Print merge_by_id
print(merge_by_id)
```

```
revenue
   branch_id    city  revenue
0         10  Austin     100
1         20  Denver      83
2         30 Springfield      4
3         47  Mendocino    200

managers
   branch_id    city  manager
0         10  Austin  Charles
1         20  Denver    Joel
2         47  Mendocino  Brett
3         31 Springfield  Sally

   branch_id_x    city  revenue  branch_id_y  manager
0         10  Austin     100         10  Charles
1         20  Denver      83         20    Joel
2         30 Springfield      4         31   Sally
3         47  Mendocino    200         47   Brett

   branch_id    city_x  revenue  city_y  manager
```

0	10	Austin	100	Austin	Charles
1	20	Denver	83	Denver	Joel
2	47	Mendocino	200	Mendocino	Brett

pd.merge seems default on inner join. Note how it handles `branch_id` with `branch_id_x/y` in the first example, and how it handles `city` with `city_x/y` in the second example. Will this happen in SQL?

Also note Notice that when you merge on 'city', the resulting DataFrame has a peculiar result: In row 2, the city Springfield has two different branch IDs. This is because there are actually two different cities named Springfield - one in the State of Illinois, and the other in Missouri. The revenue DataFrame has the one from Illinois, and the managers DataFrame has the one from Missouri. Consequently, when you merge on 'branch_id', both of these get dropped from the merged DataFrame.

Merging on columns with non-matching labels

You continue working with the revenue & managers DataFrames from before. This time, someone has changed the field name 'city' to 'branch' in the managers table. Now, when you attempt to merge DataFrames, an exception is thrown:

```
pd.merge(revenue, managers, on='city') Traceback (most recent call last): ... .. pd.merge(revenue, managers, on='city') ... .. KeyError:
'city' Given this, it will take a bit more work for you to join or merge on the city/branch name. You have to specify the left_on and right_on
parameters in the call to pd.merge().
```

As before, pandas has been pre-imported as `pd` and the revenue and managers DataFrames are in your namespace. They have been printed in the IPython Shell so you can examine the columns prior to merging.

Are you able to merge better than in the last exercise? How should the rows with Springfield be handled?

This is actually the typical SQL case

```
In [5]: import pandas as pd

revenue = pd.DataFrame({'branch_id':[10,20,30,47], 'city':['Austin','Denver','Springfield','Mendocino'],
                        'revenue':[100,83,4,200], 'state':['TX','CO','IL','CA']})
print('revenue')
print(revenue)

managers = pd.DataFrame({'branch':['Austin','Denver','Mendocino','Springfield'], 'branch_id':[10,20,47,31],
                        'manager':['Charles','Joel','Brett','Sally'], 'state':['TX','CO','CA','MO']})
print('managers')
print(managers)

# Merge revenue & managers on 'city' & 'branch': combined
combined = pd.merge(revenue, managers, left_on='city', right_on='branch')

# Print combined
print(combined)
```

```
revenue
   branch_id    city  revenue  state
0         10   Austin     100    TX
1         20   Denver      83    CO
2         30 Springfield      4    IL
3         47  Mendocino    200    CA

managers
   branch  branch_id  manager  state
0   Austin         10  Charles    TX
1   Denver         20    Joel    CO
2  Mendocino         47   Brett    CA
3  Springfield         31   Sally    MO

   branch_id_x    city  revenue  state_x    branch  branch_id_y \
0         10   Austin     100    TX    Austin         10
1         20   Denver      83    CO    Denver         20
2         30  Springfield      4    IL  Springfield         31
3         47  Mendocino    200    CA    Mendocino         47

   manager  state_y
0  Charles    TX
1    Joel    CO
2   Sally    MO
3   Brett    CA
```

Merging on multiple columns

Another strategy to disambiguate cities with identical names is to add information on the states in which the cities are located. To this end, you add a column called `state` to both DataFrames from the preceding exercises. Again, pandas has been pre-imported as `pd` and the `revenue` and `managers` DataFrames are in your namespace.

Your goal in this exercise is to use `pd.merge()` to merge DataFrames using multiple columns (using `'branch_id'`, `'city'`, and `'state'` in this case).

Are you able to match all your company's branches correctly?

INSTRUCTIONS

Create a column called `'state'` in the DataFrame `revenue`, consisting of the list `['TX','CO','IL','CA']`. Create a column called `'state'` in the DataFrame `managers`, consisting of the list `['TX','CO','CA','MO']`. Merge the DataFrames `revenue` and `managers` using three columns: `'branch_id'`, `'city'`, and `'state'`. Pass them in as a list to the `on` parameter of `pd.merge()`.

```
In [8]: import pandas as pd

revenue = pd.DataFrame({'branch_id':[10,20,30,47], 'city':['Austin','Denver','Springfield','Mendocino'],
                        'revenue':[100,83,4,200], 'state':['TX','CO','IL','CA']})

print('revenue')
print(revenue)

managers = pd.DataFrame({'branch_id':[10,20,47,31], 'city':['Austin','Denver','Mendocino','Springfield'],
                        'manager':['Charles','Joel','Brett','Sally'], 'state':['TX','CO','CA','MO']})

print('managers')
print(managers)

# Merge revenue & managers on 'branch_id', 'city', & 'state': combined
combined = pd.merge(revenue, managers, on=['branch_id', 'city', 'state'])

# Print combined
print(combined)
```

```
revenue
   branch_id    city  revenue  state
0         10   Austin     100     TX
1         20   Denver      83     CO
2         30 Springfield      4     IL
3         47  Mendocino    200     CA

managers
   branch_id    city  manager  state
0         10   Austin  Charles     TX
1         20   Denver    Joel     CO
2         47  Mendocino   Brett     CA
3         31 Springfield   Sally     MO

   branch_id    city  revenue  state  manager
0         10   Austin     100     TX   Charles
1         20   Denver      83     CO     Joel
2         47  Mendocino    200     CA     Brett
```

Joining by Index

The DataFrames `revenue` and `managers` are displayed in the IPython Shell. Here, they are indexed by 'branch_id'.

Choose the function call below that will join the DataFrames on their indexes and return 5 rows with index labels [10, 20, 30, 31, 47]. Explore each of them in the IPython Shell to get a better understanding of their functionality.

city revenue state

branch_id

10 Austin 100 TX 20 Denver 83 CO 30 Springfield 4 IL 47 Mendocino 200 CA

branch manager state

branch_id

10 Austin Charlers TX 20 Denver Joel CO 47 Mendocino Brett CA 31 Springfield Sally MO

INSTRUCTIONS

Possible Answers `pd.merge(revenue, managers, on='branch_id')`. press 1 `pd.merge(managers, revenue, how='left')`. press 2 `revenue.join(managers, lsuffix='_rev', rsuffix='_mng', how='outer')`. Answer. press 3 `managers.join(revenue, lsuffix='_mgn', rsuffix='_rev', how='left')`. press 4

Remember, the DataFrame .join() method joins on the Index while the pd.merge() function can merge on arbitrary DataFrame columns.

Choosing a joining strategy

Suppose you have two DataFrames: `students` (with columns 'StudentID', 'LastName', 'FirstName', and 'Major') and `midterm_results` (with columns 'StudentID', 'Q1', 'Q2', and 'Q3' for their scores on midterm questions).

You want to combine the DataFrames into a single DataFrame `grades`, and be able to easily spot which students wrote the midterm and which didn't (their midterm question scores 'Q1', 'Q2', & 'Q3' should be filled with NaN values).

You also want to drop rows from `midterm_results` in which the `StudentID` is not found in `students`.

Which of the following strategies gives the desired result?

INSTRUCTIONS

Possible Answers A left join: `grades = pd.merge(students, midterm_results, how='left')`. Answer press 1 A right join: `grades = pd.merge(students, midterm_results, how='right')`. press 2 An inner join: `grades = pd.merge(students, midterm_results, how='inner')`. press 3 An outer join: `grades = pd.merge(students, midterm_results, how='outer')`. press 4

Left & right merging on multiple columns

You now have, in addition to the revenue and managers DataFrames from prior exercises, a DataFrame sales that summarizes units sold from specific branches (identified by city and state but not branch_id).

Once again, the managers DataFrame uses the label branch in place of city as in the other two DataFrames. Your task here is to employ left and right merges to preserve data and identify where data is missing.

By merging revenue and sales with a right merge, you can identify the missing revenue values. Here, you don't need to specify left_on or right_on because the columns to merge on have matching labels.

By merging sales and managers with a left merge, you can identify the missing manager. Here, the columns to merge on have conflicting labels, so you must specify left_on and right_on. In both cases, you're looking to figure out how to connect the fields in rows containing Springfield.

pandas has been imported as pd and the three DataFrames revenue, managers, and sales have been pre-loaded. They have been printed for you to explore in the IPython Shell.

INSTRUCTIONS

Execute a right merge using `pd.merge()` with revenue and sales to yield a new DataFrame `revenue_and_sales`. Use `how='right'` and `on=['city', 'state']`. Print the new DataFrame `revenue_and_sales`. This has been done for you. Execute a left merge with sales and managers to yield a new DataFrame `sales_and_managers`. Use `how='left'`, `left_on=['city', 'state']`, and `right_on=['branch', 'state']`. Print the new DataFrame `sales_and_managers`. This has been done for you, so hit 'Submit Answer' to see the result!

```

In [9]: import pandas as pd

revenue = pd.DataFrame({'branch_id':[10,20,30,47], 'city':['Austin', 'Denver', 'Springfield', 'Mendocino'],
                        'revenue':[100,83,4,200], 'state':['TX', 'CO', 'IL', 'CA']})
print('revenue')
print(revenue)

managers = pd.DataFrame({'branch':['Austin', 'Denver', 'Mendocino', 'Springfield'], 'branch_id':[10,20,47,31],
                        'manager':['Charles', 'Joel', 'Brett', 'Sally'], 'state':['TX', 'CO', 'CA', 'MO']})
print('managers')
print(managers)

sales = pd.DataFrame({'city':['Mendocino', 'Denver', 'Austin', 'Springfield', 'Springfield'],
                      'state':['CA', 'CO', 'TX', 'MO', 'IL'], 'units':[1,4,2,5,1]})
print('sales')
print(sales )

#my code above

# Merge revenue and sales: revenue_and_sales
revenue_and_sales = pd.merge(revenue, sales, on=['city', 'state'], how='right')

# Print revenue_and_sales
print(revenue_and_sales)

# Merge sales and managers: sales_and_managers
sales_and_managers = pd.merge(sales, managers, left_on=['city', 'state'], right_on=['branch', 'state'], how='left')

# Print sales_and_managers
print(sales_and_managers)

```

```

revenue
   branch_id  city  revenue  state
0         10  Austin      100    TX
1         20  Denver       83    CO
2         30  Springfield     4    IL
3         47  Mendocino    200    CA

```

```

managers
   branch  branch_id  manager  state
0   Austin         10  Charles    TX
1   Denver         20    Joel    CO
2  Mendocino         47   Brett    CA

```



```

3 Springfield      31    Sally    MO
sales
      city state  units
0  Mendocino   CA     1
1    Denver   CO     4
2    Austin   TX     2
3 Springfield   MO     5
4 Springfield   IL     1
      branch_id      city  revenue  state  units
0         10.0    Austin   100.0    TX     2
1         20.0    Denver    83.0    CO     4
2         30.0 Springfield    4.0    IL     1
3         47.0  Mendocino  200.0    CA     1
4          NaN Springfield    NaN    MO     5
      city state  units      branch  branch_id  manager
0  Mendocino   CA     1    Mendocino    47.0    Brett
1    Denver   CO     4    Denver    20.0    Joel
2    Austin   TX     2    Austin    10.0  Charles
3 Springfield   MO     5 Springfield    31.0    Sally
4 Springfield   IL     1          NaN    NaN    NaN

```

Merging DataFrames with outer join

This exercise picks up where the previous one left off. The DataFrames `revenue`, `managers`, and `sales` are pre-loaded into your namespace (and, of course, pandas is imported as `pd`). Moreover, the merged DataFrames `revenue_and_sales` and `sales_and_managers` have been pre-computed exactly as you did in the previous exercise.

The merged DataFrames contain enough information to construct a DataFrame with 5 rows with all known information correctly aligned and each branch listed only once. You will try to merge the merged DataFrames on all matching keys (which computes an inner join by default). You can compare the result to an outer join and also to an outer join with restricted subset of columns as keys.

INSTRUCTIONS

Merge `sales_and_managers` with `revenue_and_sales`. Store the result as `merge_default`. Print `merge_default`. This has been done for you. Merge `sales_and_managers` with `revenue_and_sales` using `how='outer'`. Store the result as `merge_outer`. Print `merge_outer`. This has been done for you. Merge `sales_and_managers` with `revenue_and_sales` only on `['city','state']` using an outer join. Store the result as `merge_outer_on` and hit 'Submit Answer' to see what the merged DataFrames look like!

It seems: outer here means full outer join but not include other outer joins such as left and right outer joins. default is inner join

Also note, all outer joins, including full outer join can be used with ON conditions. Usually this will reduce the number of rows.

```

In [10]: #Need data from previous cell
# Perform the first merge: merge_default
merge_default = pd.merge(sales_and_managers, revenue_and_sales)

# Print merge_default
print(merge_default)

# Perform the second merge: merge_outer
merge_outer = pd.merge(sales_and_managers, revenue_and_sales, how='outer')

# Print merge_outer
print(merge_outer)

# Perform the third merge: merge_outer_on
merge_outer_on = pd.merge(sales_and_managers, revenue_and_sales, on=['city', 'state'], how='outer')

# Print merge_outer_on
print(merge_outer_on)

```

	city	state	units	branch	branch_id	manager	revenue
0	Mendocino	CA	1	Mendocino	47.0	Brett	200.0
1	Denver	CO	4	Denver	20.0	Joel	83.0
2	Austin	TX	2	Austin	10.0	Charles	100.0

	city	state	units	branch	branch_id	manager	revenue
0	Mendocino	CA	1	Mendocino	47.0	Brett	200.0
1	Denver	CO	4	Denver	20.0	Joel	83.0
2	Austin	TX	2	Austin	10.0	Charles	100.0
3	Springfield	MO	5	Springfield	31.0	Sally	NaN
4	Springfield	IL	1	NaN	NaN	NaN	NaN
5	Springfield	IL	1	NaN	30.0	NaN	4.0
6	Springfield	MO	5	NaN	NaN	NaN	NaN

	city	state	units_x	branch	branch_id_x	manager	branch_id_y	\
0	Mendocino	CA	1	Mendocino	47.0	Brett	47.0	
1	Denver	CO	4	Denver	20.0	Joel	20.0	
2	Austin	TX	2	Austin	10.0	Charles	10.0	
3	Springfield	MO	5	Springfield	31.0	Sally	NaN	
4	Springfield	IL	1	NaN	NaN	NaN	30.0	

	revenue	units_y
0	200.0	1
1	83.0	4

2	100.0	2
3	NaN	5
4	4.0	1

Using merge_ordered()

This exercise uses pre-loaded DataFrames `austin` and `houston` that contain weather data from the cities Austin and Houston respectively. They have been printed in the IPython Shell for you to examine.

Weather conditions were recorded on separate days and you need to merge these two DataFrames together such that the dates are ordered. To do this, you'll use `pd.merge_ordered()`. After you're done, note the order of the rows before and after merging.

INSTRUCTIONS

Perform an ordered merge on `austin` and `houston` using `pd.merge_ordered()`. Store the result as `tx_weather`. Print `tx_weather`. You should notice that the rows are sorted by the date but it is not possible to tell which observation came from which city. Perform another ordered merge on `austin` and `houston`. This time, specify the keyword arguments `on='date'` and `suffixes=['_aus', '_hus']` so that the rows can be distinguished. Store the result as `tx_weather_suff`. Print `tx_weather_suff` to examine its contents. This has been done for you. Perform a third ordered merge on `austin` and `houston`. This time, in addition to the `on` and `suffixes` parameters, specify the keyword argument `fill_method='ffill'` to use forward-filling to replace NaN entries with the most recent non-null entry, and hit 'Submit Answer' to examine the contents of the merged DataFrames!

```
In [18]: import pandas as pd
austin = pd.DataFrame({'date': ['2016-01-01', '2016-02-08', '2016-01-17'], 'ratings': ['Cloudy', 'Cloudy', 'Sunny']})
austin['date'] = pd.to_datetime(austin['date'])
houston = pd.DataFrame({'date': ['2016-01-04', '2016-01-01', '2016-03-01'], 'ratings': ['Rainy', 'Cloudy', 'Sunny']})
houston['date'] = pd.to_datetime(houston['date'])

#my code above

# Perform the first ordered merge: tx_weather
tx_weather = pd.merge_ordered(austin, houston)

# Print tx_weather
print(tx_weather)

# Perform the second ordered merge: tx_weather_suff
tx_weather_suff = pd.merge_ordered(austin, houston, on='date', suffixes=['_aus', '_hus'])

# Print tx_weather_suff
print(tx_weather_suff)

# Perform the third ordered merge: tx_weather_ffill
tx_weather_ffill = pd.merge_ordered(austin, houston, on='date', fill_method='ffill', suffixes=['_aus', '_hus'])

# Print tx_weather_ffill
print(tx_weather_ffill)
```

```

      date ratings
0 2016-01-01  Cloudy
1 2016-01-04   Rainy
2 2016-01-17   Sunny
3 2016-02-08  Cloudy
4 2016-03-01   Sunny
      date_aus ratings  date_hus
0 2016-01-01  Cloudy 2016-01-01
1 2016-02-08  Cloudy 2016-01-01
2         NaT   Rainy 2016-01-04
3 2016-01-17   Sunny 2016-03-01
      date ratings_aus ratings_hus
0 2016-01-01     Cloudy     Cloudy
1 2016-01-04     Cloudy     Rainy
2 2016-01-17     Sunny     Rainy
```

3	2016-02-08	Cloudy	Rainy
4	2016-03-01	Cloudy	Sunny

We can also order by other columns

However, note the `merger_ordered()` here has nothing to do with the `order by` in `sql`. In `pandas`, I think `sort` is related to `order by`.

Using `merge_asof()`

Similar to `pd.merge_ordered()`, the `pd.merge_asof()` function will also merge values in order using the `on` column, but for each row in the left DataFrame, only rows from the right DataFrame whose `'on'` column values are less than the left value will be kept.

This function can be used to align disparate datetime frequencies without having to first resample.

Here, you'll merge monthly oil prices (US dollars) into a full automobile fuel efficiency dataset. The oil and automobile DataFrames have been pre-loaded as `oil` and `auto`. The first 5 rows of each have been printed in the IPython Shell for you to explore.

These datasets will align such that the first price of the year will be broadcast into the rows of the automobiles DataFrame. This is considered correct since by the start of any given year, most automobiles for that year will have already been manufactured.

You'll then inspect the merged DataFrame, resample by year and compute the mean `'Price'` and `'mpg'`. You should be able to see a trend in these two columns, that you can confirm by computing the Pearson correlation between resampled `'Price'` and `'mpg'`.

INSTRUCTIONS

Merge `auto` and `oil` using `pd.merge_asof()` with `left_on='yr'` and `right_on='Date'`. Store the result as `merged`. Print the tail of `merged`. This has been done for you. Resample `merged` using `'A'` (annual frequency), and `on='Date'`. Select `[['mpg','Price']]` and aggregate the mean. Store the result as `yearly`. Hit Submit Answer to examine the contents of `yearly` and `yearly.corr()`, which shows the Pearson correlation between the resampled `'Price'` and `'mpg'`.

```
In [35]: import pandas as pd
auto = pd.read_csv("automobiles.csv")
auto['yr'] = pd.to_datetime(auto['yr'])
oil = pd.read_csv("oil_price.csv")
oil['Date'] = pd.to_datetime(oil['Date'])
print(auto.head())
print('-----')
print(oil.head())
print('-----')
#above my code

# Merge auto and oil: merged
merged = pd.merge_asof(auto, oil, left_on='yr', right_on='Date')

# Print the tail of merged
print(merged.tail())
print('-----')

# Resample merged: yearly
yearly = merged.resample('A', on='Date')[['mpg', 'Price']].mean()

# Print yearly
print(yearly)
print('-----')

# Print yearly.corr()
print(yearly.corr())
print('-----')
```

	mpg	cyl	displ	hp	weight	accel	yr	origin	\
0	18.0	8	307.0	130	3504	12.0	1970-01-01	US	
1	15.0	8	350.0	165	3693	11.5	1970-01-01	US	
2	18.0	8	318.0	150	3436	11.0	1970-01-01	US	
3	16.0	8	304.0	150	3433	12.0	1970-01-01	US	
4	17.0	8	302.0	140	3449	10.5	1970-01-01	US	

	name
0	chevrolet chevelle malibu
1	buick skylark 320
2	plymouth satellite
3	amc rebel sst
4	ford torino

```

-----
      Date  Price
0 1970-01-01    3.35
1 1970-02-01    3.35
2 1970-03-01    3.35
3 1970-04-01    3.35
4 1970-05-01    3.35
-----

```

```

      mpg  cyl  displ  hp  weight  accel      yr  origin      name \
387  27.0    4   140.0  86   2790   15.6 1982-01-01    US  ford mustang gl
388  44.0    4    97.0  52   2130   24.6 1982-01-01  Europe    vw pickup
389  32.0    4   135.0  84   2295   11.6 1982-01-01    US   dodge rampage
390  28.0    4   120.0  79   2625   18.6 1982-01-01    US   ford ranger
391  31.0    4   119.0  82   2720   19.4 1982-01-01    US   chevy s-10

```

```

      Date  Price
387 1982-01-01  33.85
388 1982-01-01  33.85
389 1982-01-01  33.85
390 1982-01-01  33.85
391 1982-01-01  33.85
-----

```

```

      mpg  Price
Date
1970-12-31  17.689655    3.35
1971-12-31  21.111111    3.56
1972-12-31  18.714286    3.56
1973-12-31  17.100000    3.56
1974-12-31  22.769231   10.11
1975-12-31  20.266667   11.16
1976-12-31  21.573529   11.16
1977-12-31  23.375000   13.90
1978-12-31  24.061111   14.85
1979-12-31  25.093103   14.85
1980-12-31  33.803704   32.50
1981-12-31  30.185714   38.00
1982-12-31  32.000000   33.85
-----

```

```

      mpg      Price
mpg    1.000000  0.948677
Price  0.948677  1.000000
-----

```


Case Study - Summer Olympics

To cement your new skills, you'll apply them by working on an in-depth study involving Olympic medal data. The analysis involves integrating your multi-DataFrame skills from this course and also skills you've gained in previous pandas courses. This is a rich dataset that will allow you to fully leverage your pandas data manipulation skills. Enjoy!

Loading Olympic edition DataFrame

In this chapter, you'll be using The Guardian's Olympic medal dataset.

Your first task here is to prepare a DataFrame editions from a tab-separated values (TSV) file.

Initially, editions has 26 rows (one for each Olympic edition, i.e., a year in which the Olympics was held) and 7 columns: 'Edition', 'Bronze', 'Gold', 'Silver', 'Grand Total', 'City', and 'Country'.

For the analysis that follows, you won't need the overall medal counts, so you want to keep only the useful columns from editions: 'Edition', 'Grand Total', City, and Country.

INSTRUCTIONS

Read `file_path` into a DataFrame called `editions`. The identifier `file_path` has been pre-defined with the filename 'Summer Olympic medallists 1896 to 2008 - EDITIONS.tsv'. You'll have to use the option `sep='\t'` because the file uses tabs to delimit fields (`pd.read_csv()` expects commas by default). Select only the columns 'Edition', 'Grand Total', 'City', and 'Country' from `editions`. Print the final DataFrame `editions` in entirety (there are only 26 rows). This has been done for you, so hit 'Submit Answer' to see the result!

Note .tsv is just tab separated value file

```
In [38]: # Import pandas
import pandas as pd

# Load DataFrame from file_path: editions
editions = pd.read_csv('Summer Olympic medalists 1896 to 2008 - EDITIONS.tsv', sep='\t')

# Extract the relevant columns: editions
editions = editions[['Edition', 'Grand Total', 'City', 'Country']]

# Print editions DataFrame
print(editions)
```

	Edition	Grand Total	City	Country
0	1896	151	Athens	Greece
1	1900	512	Paris	France
2	1904	470	St. Louis	United States
3	1908	804	London	United Kingdom
4	1912	885	Stockholm	Sweden
5	1920	1298	Antwerp	Belgium
6	1924	884	Paris	France
7	1928	710	Amsterdam	Netherlands
8	1932	615	Los Angeles	United States
9	1936	875	Berlin	Germany
10	1948	814	London	United Kingdom
11	1952	889	Helsinki	Finland
12	1956	885	Melbourne	Australia
13	1960	882	Rome	Italy
14	1964	1010	Tokyo	Japan
15	1968	1031	Mexico City	Mexico
16	1972	1185	Munich	West Germany (now Germany)
17	1976	1305	Montreal	Canada
18	1980	1387	Moscow	U.S.S.R. (now Russia)
19	1984	1459	Los Angeles	United States
20	1988	1546	Seoul	South Korea
21	1992	1705	Barcelona	Spain
22	1996	1859	Atlanta	United States
23	2000	2015	Sydney	Australia
24	2004	1998	Athens	Greece
25	2008	2042	Beijing	China

Loading IOC codes DataFrame

Your task here is to prepare a DataFrame `ioc_codes` from a comma-separated values (CSV) file.

Initially, `ioc_codes` has 200 rows (one for each country) and 3 columns: 'Country', 'NOC', & 'ISO code'.

For the analysis that follows, you want to keep only the useful columns from `ioc_codes`: 'Country' and 'NOC' (the column 'NOC' contains three-letter codes representing each country).

INSTRUCTIONS

Read `file_path` into a DataFrame called `ioc_codes`. The identifier `file_path` has been pre-defined with the filename 'Summer Olympic medalists 1896 to 2008 - IOC COUNTRY CODES.csv'. Select only the columns 'Country' and 'NOC' from `ioc_codes`. Print the leading 5 and trailing 5 rows of the DataFrame `ioc_codes` (there are 200 rows in total). This has been done for you, so hit 'Submit Answer' to see the result!

```
In [40]: # Import pandas
import pandas as pd

# Load DataFrame from file_path: ioc_codes
ioc_codes = pd.read_csv('Summer Olympic medalists 1896 to 2008 - IOC COUNTRY CODES.csv')

# Extract the relevant columns: ioc_codes
ioc_codes = ioc_codes[['Country', 'NOC']]

# Print first and last 5 rows of ioc_codes
print(ioc_codes.head())
print(ioc_codes.tail())
```

	Country	NOC
0	Afghanistan	AFG
1	Albania	ALB
2	Algeria	ALG
3	American Samoa*	ASA
4	Andorra	AND
	Country	NOC
196	Vietnam	VIE
197	Virgin Islands*	ISV
198	Yemen	YEM
199	Zambia	ZAM
200	Zimbabwe	ZIM

Building medals DataFrame

Here, you'll start with the DataFrame editions from the previous exercise.

You have a sequence of files `summer_1896.csv`, `summer_1900.csv`, ..., `summer_2008.csv`, one for each Olympic edition (year).

You will build up a dictionary `medals_dict` with the Olympic editions (years) as keys and DataFrames as values.

The dictionary is built up inside a loop over the year of each Olympic edition (from the Index of editions).

Once the dictionary of DataFrames is built up, you will combine the DataFrames using `pd.concat()`.

INSTRUCTIONS

Within the for loop: Create the file path. This has been done for you. Read `file_path` into a DataFrame. Assign the result to the year key of `medals_dict`. Select only the columns 'Athlete', 'NOC', and 'Medal' from `medals_dict[year]`. Create a new column called 'Edition' in the DataFrame `medals_dict[year]` whose entries are all year. Concatenate the dictionary of DataFrames `medals_dict` into a DataFrame called `medals`. Specify the keyword argument `ignore_index=True` to prevent repeated integer indices. Print the first and last 5 rows of `medals`. This has been done for you, so hit 'Submit Answer' to see the result!

Note for .tsv file, I need `sep = '\t'` option

```
In [8]: # Import pandas
import pandas as pd

##### Need run the following code to create files if not existed.

# editions = pd.read_csv('Summer Olympic medalists 1896 to 2008 - EDITIONS.tsv', sep='\t')
# # Extract the relevant columns: editions
# editions = editions[['Edition', 'Grand Total', 'City', 'Country']]

# df = pd.read_csv("Summer Olympic medalists 1896 to 2008.tsv", sep='\t')

# for year in editions['Edition']:
#     df_temp = df[df['Edition']== year]
#     file_path = 'summer_{:d}.csv'.format(year)
#     df_temp.to_csv(file_path)

#my code above

# Create empty dictionary: medals_dict
medals_dict = {}

for year in editions['Edition']:

    # Create the file path: file_path
    file_path = 'summer_{:d}.csv'.format(year)

    # Load file_path into a DataFrame: medals_dict[year]
    medals_dict[year] = pd.read_csv(file_path)

    # Extract relevant columns: medals_dict[year]
    medals_dict[year] = medals_dict[year][['Athlete', 'NOC', 'Medal']]

    # Assign year to column 'Edition' of medals_dict
    medals_dict[year]['Edition'] = year

# Concatenate medals_dict: medals
medals = pd.concat(medals_dict, ignore_index=True)

# Print first and last 5 rows of medals
print(medals.head())
print(medals.tail())
```

	Athlete	NOC	Medal	Edition
0	HAJOS, Alfred	HUN	Gold	1896
1	HERSCHMANN, Otto	AUT	Silver	1896
2	DRIVAS, Dimitrios	GRE	Bronze	1896
3	MALOKINIS, Ioannis	GRE	Gold	1896
4	CHASAPIS, Spiridon	GRE	Silver	1896
	Athlete	NOC	Medal	Edition
29211	ENGLISH, Mirko	GER	Silver	2008
29212	MIZGAITIS, Mindaugas	LTU	Bronze	2008
29213	PATRIKEEV, Yuri	ARM	Bronze	2008
29214	LOPEZ, Mijain	CUB	Gold	2008
29215	BAROEV, Khasan	RUS	Silver	2008

Counting medals by country/edition in a pivot table

Here, you'll start with the concatenated DataFrame medals from the previous exercise.

You can construct a pivot table to see the number of medals each country won in each year. The result is a new DataFrame with the Olympic edition on the Index and with 138 country NOC codes as columns. If you want a refresher on pivot tables, it may be useful to refer back to the relevant exercises in Manipulating DataFrames with pandas.

INSTRUCTIONS

Construct a pivot table from the DataFrame medals, aggregating by count (by specifying the aggfunc parameter). Use 'Edition' as the index, 'Athlete' for the values, and 'NOC' for the columns. Print the first & last 5 rows of medal_counts. This has been done for you, so hit 'Submit Answer' to see the results!

```
In [9]: # Construct the pivot_table: medal_counts
medal_counts = medals.pivot_table(index='Edition', values='Athlete', columns='NOC', aggfunc='count')

# Print the first & last 5 rows of medal_counts
print(medal_counts.head())
print(medal_counts.tail())
```

NOC	AFG	AHO	ALG	ANZ	ARG	ARM	AUS	AUT	AZE	BAH	...	URS	URU	\
Edition											...			
1896	NaN	NaN	NaN	NaN	NaN	NaN	2.0	5.0	NaN	NaN	...	NaN	NaN	
1900	NaN	NaN	NaN	NaN	NaN	NaN	5.0	6.0	NaN	NaN	...	NaN	NaN	
1904	NaN	NaN	NaN	NaN	NaN	NaN	NaN	1.0	NaN	NaN	...	NaN	NaN	
1908	NaN	NaN	NaN	19.0	NaN	NaN	NaN	1.0	NaN	NaN	...	NaN	NaN	
1912	NaN	NaN	NaN	10.0	NaN	NaN	NaN	14.0	NaN	NaN	...	NaN	NaN	

NOC	USA	UZB	VEN	VIE	YUG	ZAM	ZIM	ZZX
Edition								
1896	20.0	NaN	NaN	NaN	NaN	NaN	NaN	6.0
1900	55.0	NaN	NaN	NaN	NaN	NaN	NaN	34.0
1904	394.0	NaN	NaN	NaN	NaN	NaN	NaN	8.0
1908	63.0	NaN	NaN	NaN	NaN	NaN	NaN	NaN
1912	101.0	NaN	NaN	NaN	NaN	NaN	NaN	NaN

[5 rows x 138 columns]

NOC	AFG	AHO	ALG	ANZ	ARG	ARM	AUS	AUT	AZE	BAH	...	URS	URU	\
Edition											...			
1992	NaN	NaN	2.0	NaN	2.0	NaN	57.0	6.0	NaN	1.0	...	NaN	NaN	
1996	NaN	NaN	3.0	NaN	20.0	2.0	132.0	3.0	1.0	5.0	...	NaN	NaN	
2000	NaN	NaN	5.0	NaN	20.0	1.0	183.0	4.0	3.0	6.0	...	NaN	1.0	
2004	NaN	NaN	NaN	NaN	47.0	NaN	157.0	8.0	5.0	2.0	...	NaN	NaN	
2008	1.0	NaN	2.0	NaN	51.0	6.0	149.0	3.0	7.0	5.0	...	NaN	NaN	

NOC	USA	UZB	VEN	VIE	YUG	ZAM	ZIM	ZZX
Edition								
1992	224.0	NaN	NaN	NaN	NaN	NaN	NaN	NaN
1996	260.0	2.0	NaN	NaN	26.0	1.0	NaN	NaN
2000	248.0	4.0	NaN	1.0	26.0	NaN	NaN	NaN
2004	264.0	5.0	2.0	NaN	NaN	NaN	3.0	NaN
2008	315.0	6.0	1.0	1.0	NaN	NaN	4.0	NaN

[5 rows x 138 columns]

Computing fraction of medals per Olympic edition

In this exercise, you'll start with the DataFrames `editions`, `medals`, & `medal_counts` from prior exercises.

You can extract a Series with the total number of medals awarded in each Olympic edition.

The DataFrame `medal_counts` can be divided row-wise by the total number of medals awarded each edition; the method `.divide()` performs the broadcast as you require.

This gives you a normalized indication of each country's performance in each edition.

INSTRUCTIONS

Set the index of the DataFrame `editions` to be 'Edition' (using the method `.set_index()`). Save the result as `totals`. Extract the 'Grand Total' column from `totals` and assign the result back to `totals`. Divide the DataFrame `medal_counts` by `totals` along each row. You will have to use the `.divide()` method with the option `axis='rows'`. Assign the result to `fractions`. Print first & last 5 rows of the DataFrame `fractions`. This has been done for you, so hit 'Submit Answer' to see the results!


```
In [10]: # Set Index of editions: totals
totals = editions.set_index('Edition')

# Reassign totals['Grand Total']: totals
totals = totals['Grand Total']

# Divide medal_counts by totals: fractions
fractions = medal_counts.divide(totals, axis='rows')

# Print first & last 5 rows of fractions
print(fractions.head())
print(fractions.tail())
```

NOC	AFG	AHO	ALG	ANZ	ARG	ARM	AUS	AUT	AZE	BAH	\
Edition											
1896	NaN	NaN	NaN	NaN	NaN	NaN	0.013245	0.033113	NaN	NaN	
1900	NaN	NaN	NaN	NaN	NaN	NaN	0.009766	0.011719	NaN	NaN	
1904	NaN	NaN	NaN	NaN	NaN	NaN	NaN	0.002128	NaN	NaN	
1908	NaN	NaN	NaN	0.023632	NaN	NaN	NaN	0.001244	NaN	NaN	
1912	NaN	NaN	NaN	0.011299	NaN	NaN	NaN	0.015819	NaN	NaN	

NOC	...	URS	URU	USA	UZB	VEN	VIE	YUG	ZAM	ZIM	ZZX
Edition	...										
1896	...	NaN	NaN	0.132450	NaN	NaN	NaN	NaN	NaN	NaN	0.039735
1900	...	NaN	NaN	0.107422	NaN	NaN	NaN	NaN	NaN	NaN	0.066406
1904	...	NaN	NaN	0.838298	NaN	NaN	NaN	NaN	NaN	NaN	0.017021
1908	...	NaN	NaN	0.078358	NaN	NaN	NaN	NaN	NaN	NaN	NaN
1912	...	NaN	NaN	0.114124	NaN	NaN	NaN	NaN	NaN	NaN	NaN

[5 rows x 138 columns]

NOC	AFG	AHO	ALG	ANZ	ARG	ARM	AUS	AUT	\
Edition									
1992	NaN	NaN	0.001173	NaN	0.001173	NaN	0.033431	0.003519	
1996	NaN	NaN	0.001614	NaN	0.010758	0.001076	0.071006	0.001614	
2000	NaN	NaN	0.002481	NaN	0.009926	0.000496	0.090819	0.001985	
2004	NaN	NaN	NaN	NaN	0.023524	NaN	0.078579	0.004004	
2008	0.00049	NaN	0.000979	NaN	0.024976	0.002938	0.072968	0.001469	

NOC	AZE	BAH	...	URS	URU	USA	UZB	VEN	\
Edition			...						
1992	NaN	0.000587	...	NaN	NaN	0.131378	NaN	NaN	
1996	0.000538	0.002690	...	NaN	NaN	0.139860	0.001076	NaN	

2000	0.001489	0.002978	...	NaN	0.000496	0.123077	0.001985	NaN
2004	0.002503	0.001001	...	NaN	NaN	0.132132	0.002503	0.001001
2008	0.003428	0.002449	...	NaN	NaN	0.154261	0.002938	0.000490
NOC	VIE	YUG	ZAM	ZIM	ZZX			
Edition								
1992	NaN	NaN	NaN	NaN	NaN			
1996	NaN	0.013986	0.000538	NaN	NaN			
2000	0.000496	0.012903	NaN	NaN	NaN			
2004	NaN	NaN	NaN	0.001502	NaN			
2008	0.000490	NaN	NaN	0.001959	NaN			

[5 rows x 138 columns]

We have used `.multiply()` before and here we have `.divide()`

Computing percentage change in fraction of medals won

Here, you'll start with the DataFrames `editions`, `medals`, `medal_counts`, & `fractions` from prior exercises.

To see if there is a host country advantage, you first want to see how the fraction of medals won changes from edition to edition.

The expanding mean provides a way to see this down each column. It is the value of the mean with all the data available up to that point in time. If you are interested in learning more about pandas' expanding transformations, this section of the pandas documentation has additional information.

INSTRUCTIONS

Create `mean_fractions` by chaining the methods `.expanding().mean()` to `fractions`. Compute the percentage change in `mean_fractions` down each column by applying `.pct_change()` and multiplying by 100. Assign the result to `fractions_change`. Reset the index of `fractions_change` using the `.reset_index()` method. This will make 'Edition' an ordinary column. Print the first and last 5 rows of the DataFrame `fractions_change`. This has been done for you, so hit 'Submit Answer' to see the results!

```
In [11]: # Apply the expanding mean: mean_fractions
mean_fractions = fractions.expanding().mean()

# Compute the percentage change: fractions_change
fractions_change = mean_fractions.pct_change()*100

# Reset the index of fractions_change: fractions_change
fractions_change = fractions_change.reset_index()

# Print first & last 5 rows of fractions_change
print(fractions_change.head())
print(fractions_change.tail())
```

NOC	Edition	AFG	AHO	ALG	ANZ	ARG	ARM	AUS	AUT	AZE	\
0	1896	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
1	1900	NaN	NaN	NaN	NaN	NaN	NaN	-13.134766	-32.304688	NaN	
2	1904	NaN	NaN	NaN	NaN	NaN	NaN	0.000000	-30.169386	NaN	
3	1908	NaN	NaN	NaN	NaN	NaN	NaN	0.000000	-23.013510	NaN	
4	1912	NaN	NaN	NaN	-26.092774	NaN	NaN	0.000000	6.254438	NaN	

NOC	...	URS	URU	USA	UZB	VEN	VIE	YUG	ZAM	ZIM	ZZX
0	...	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
1	...	NaN	NaN	-9.448242	NaN	NaN	NaN	NaN	NaN	NaN	33.561198
2	...	NaN	NaN	199.651245	NaN	NaN	NaN	NaN	NaN	NaN	-22.642384
3	...	NaN	NaN	-19.549222	NaN	NaN	NaN	NaN	NaN	NaN	0.000000
4	...	NaN	NaN	-12.105733	NaN	NaN	NaN	NaN	NaN	NaN	0.000000

[5 rows x 139 columns]

NOC	Edition	AFG	AHO	ALG	ANZ	ARG	ARM	AUS	\
21	1992	NaN	0.0	-7.214076	0.0	-6.767308	NaN	2.754114	
22	1996	NaN	0.0	8.959211	0.0	1.306696	NaN	10.743275	
23	2000	NaN	0.0	19.762488	0.0	0.515190	-26.935484	12.554986	
24	2004	NaN	0.0	0.000000	0.0	9.625365	0.000000	8.161162	
25	2008	NaN	0.0	-8.197807	0.0	8.588555	91.266408	6.086870	

NOC	AUT	AZE	...	URS	URU	USA	UZB	VEN	\
21	-3.034840	NaN	...	0.0	0.000000	-1.329330	NaN	0.000000	
22	-3.876773	NaN	...	0.0	0.000000	-1.010378	NaN	0.000000	
23	-3.464221	88.387097	...	0.0	-12.025323	-1.341842	42.258065	0.000000	
24	-2.186922	48.982144	...	0.0	0.000000	-1.031922	21.170339	-1.615969	
25	-3.389836	31.764436	...	0.0	0.000000	-0.450031	14.610625	-6.987342	

NOC	VIE	YUG	ZAM	ZIM	ZZX
21	NaN	0.000000	0.000000	0.000000	0.0
22	NaN	-2.667732	-10.758472	0.000000	0.0
23	NaN	-2.696445	0.000000	0.000000	0.0
24	0.000000	0.000000	0.000000	-43.491929	0.0
25	-0.661117	0.000000	0.000000	-23.316533	0.0

[5 rows x 139 columns]

Building hosts DataFrame

Your task here is to prepare a DataFrame hosts by left joining editions and ioc_codes.

Once created, you will subset the Edition and NOC columns and set Edition as the Index.

There are some missing NOC values; you will set those explicitly.

Finally, you'll reset the Index & print the final DataFrame.

INSTRUCTIONS

Create the DataFrame hosts by doing a left join on DataFrames editions and ioc_codes (using `pd.merge()`). Clean up hosts by subsetting and setting the Index. Extract the columns 'Edition' and 'NOC'. Set 'Edition' column as the Index. Use the `.loc[]` accessor to find and assign the missing values to the 'NOC' column in hosts. This has been done for you. Reset the index of hosts using `.reset_index()`, which returns a new DataFrame. Hit 'Submit Answer' to see what hosts looks like!

```

In [13]: # Import pandas
import pandas as pd

# Load DataFrame from file_path: ioc_codes
ioc_codes = pd.read_csv('Summer Olympic medalists 1896 to 2008 - IOC COUNTRY CODES.csv')

# Extract the relevant columns: ioc_codes
ioc_codes = ioc_codes[['Country', 'NOC']]

# my code above

# Left join editions and ioc_codes: hosts
hosts = pd.merge(editions, ioc_codes, how='left')

# Extract relevant columns and set index: hosts
hosts = hosts[['Edition', 'NOC']].set_index('Edition')

# Fix missing 'NOC' values of hosts
print(hosts.loc[hosts.NOC.isnull()])
hosts.loc[1972, 'NOC'] = 'FRG'
hosts.loc[1980, 'NOC'] = 'URS'
hosts.loc[1988, 'NOC'] = 'KOR'

# Reset Index of hosts: hosts
hosts = hosts.reset_index()

# Print hosts
print(hosts)

```

		NOC
Edition		
1972		NaN
1980		NaN
1988		NaN
	Edition	NOC
0	1896	GRE
1	1900	FRA
2	1904	USA
3	1908	GBR
4	1912	SWE
5	1920	BEL
6	1924	FRA

7	1928	NED
8	1932	USA
9	1936	GER
10	1948	GBR
11	1952	FIN
12	1956	AUS
13	1960	ITA
14	1964	JPN
15	1968	MEX
16	1972	FRG
17	1976	CAN
18	1980	URS
19	1984	USA
20	1988	KOR
21	1992	ESP
22	1996	USA
23	2000	AUS
24	2004	GRE
25	2008	CHN

Reshaping for analysis

This exercise starts off with `fractions_change` and `hosts` already loaded.

Your task here is to reshape the `fractions_change` DataFrame for later analysis.

Initially, `fractions_change` is a wide DataFrame of 26 rows (one for each Olympic edition) and 139 columns (one for the edition and 138 for the competing countries).

On reshaping with `pd.melt()`, as you will see, the result is a tall DataFrame with 3588 rows and 3 columns that summarizes the fractional change in the expanding mean of the percentage of medals won for each country in blocks.

INSTRUCTIONS 100 XP Create a DataFrame reshaped by reshaping the DataFrame `fractions_change` with `pd.melt()`. You'll need to use the keyword argument `id_vars='Edition'` to set the identifier variable. You'll also need to use the keyword argument `value_name='Change'` to set the measured variables. Print the shape of the DataFrames reshaped and `fractions_change`. This has been done for you. Create a DataFrame `chn` by extracting all the rows from reshaped in which the three letter code for each country ('NOC') is 'CHN'. Print the last 5 rows of the DataFrame `chn` using the `.tail()` method. This has been done for you, so hit 'Submit Answer' to see the results!

```
In [14]: # Import pandas
import pandas as pd

# Reshape fractions_change: reshaped
reshaped = pd.melt(fractions_change, id_vars='Edition', value_name='Change')

# Print reshaped.shape and fractions_change.shape
print(reshaped.shape, fractions_change.shape)

# Extract rows from reshaped where 'NOC' == 'CHN': chn
chn = reshaped.loc[reshaped.NOC == 'CHN']

# Print last 5 rows of chn
print(chn.tail())
```

```
(3588, 3) (26, 139)
   Edition  NOC  Change
567    1992  CHN   4.240630
568    1996  CHN   7.860247
569    2000  CHN  -3.851278
570    2004  CHN   0.128863
571    2008  CHN  13.251332
```

On looking at the hosting countries from the last 5 Olympic editions and the fractional change of medals won by China the last 5 editions, you can see that China fared significantly better in 2008 (i.e., when China was the host country).

Merging to compute influence

This exercise starts off with the DataFrames `reshaped` and `hosts` in the namespace.

Your task is to merge the two DataFrames and tidy the result.

The end result is a DataFrame summarizing the fractional change in the expanding mean of the percentage of medals won for the host country in each Olympic edition.

INSTRUCTIONS

Merge reshaped and hosts using an inner join. Remember, how='inner' is the default behavior for pd.merge(). Print the first 5 rows of the DataFrame merged. This has been done for you. You should see that the rows are jumbled chronologically. Set the index of merged to be 'Edition' and sort the index. Print the first 5 rows of the DataFrame influence. This has been done for you, so hit 'Submit Answer' to see the results!

```
In [15]: # Import pandas
import pandas as pd

# Merge reshaped and hosts: merged
merged = pd.merge(reshaped, hosts)

# Print first 5 rows of merged
print(merged.head())

# Set Index of merged and sort it: influence
influence = merged.set_index('Edition').sort_index()

# Print first 5 rows of influence
print(influence.head())
```

	Edition	NOC	Change
0	1956	AUS	54.615063
1	2000	AUS	12.554986
2	1920	BEL	54.757887
3	1976	CAN	-2.143977
4	2008	CHN	13.251332

	NOC	Change
Edition		
1896	GRE	NaN
1900	FRA	198.002486
1904	USA	199.651245
1908	GBR	134.489218
1912	SWE	71.896226

Plotting influence of host country

This final exercise starts off with the DataFrames influence and editions in the namespace. Your job is to plot the influence of being a host country.

INSTRUCTIONS

Create a Series called change by extracting the 'Change' column from influence. Create a bar plot of change using the .plot() method with kind='bar'. Save the result as ax to permit further customization. Customize the bar plot of change to improve readability: Apply the method .set_ylabel("% Change of Host Country Medal Count") to ax. Apply the method .set_title("Is there a Host Country Advantage?") to ax. Apply the method .set_xticklabels(editions['City']) to ax. Reveal the final plot using plt.show().

```
In [17]: # Import pyplot
import matplotlib.pyplot as plt

# Extract influence['Change']: change
change = influence['Change']

# Make bar plot of change: ax
ax = change.plot(kind='bar')

# Customize the plot to improve readability
ax.set_ylabel("% Change of Host Country Medal Count")
ax.set_title("Is there a Host Country Advantage?")
ax.set_xticklabels(editions['City'])

# Display the plot
plt.show()
```



